

# Programming Interview Practice

Problems taken from “Programming Interviews Exposed: Secrets to Landing Your Next Job” by John Mongan and Noah Suojanen

Repository of Materials: [GitHub Repo](#)

Language of Choice: Python

## Chapter 3: Linked Lists

**Problem:** *Stack Implementation.* Implement a stack using either a linked list or dynamic array (justify your decision). Design the interface to your stack to be complete, consistent, and easy to use

### Solution

I choose to implement the stack as a linked list. The disadvantage of a linked list is that you need to iterate through the list to find an element, however, since we are implementing a stack, we only need to access the first element. Dynamic arrays have convenient, index-based element access for interacting with the top of the stack, in practice will be less efficient for a stack. The reasoning here is that arrays store their data in contiguous memory locations, so every time elements are popped or pushed onto the stack, the memory addresses of all elements will need to be updated, making these operations  $O(n)$ . A linked list will only require the adjustment of a couple pointers, which will be  $O(1)$  complexity.

```
class DoublyLinkedListNode:
    def __init__(self, data):
        self.data = data
        self.next = None
        self.prev = None

class Stack:
    def __init__(self):
        self.head = None

    def push(self, data):
        new_node = DoublyLinkedListNode(data)
        if self.head is None:
            self.head = new_node
        else:
            new_node.next = self.head
            self.head.prev = new_node
            self.head = new_node
```

```
def pop(self):
    if head is None:
        raise IndexError("Cannot pop from an empty Stack")
    data_to_pop = self.head.data
    self.head = self.head.next
    self.head.prev = None
    return data_to_pop

def peek:
    if head is None:
        raise IndexError("Cannot peek from an empty Stack")
    return self.head.data
```

Notes based on book solution:

- A dynamic array doesn't require additional memory allocation for pointers, so if the elements are all small integers (require much less memory than a pointer) then a dynamic array may use less memory space
- LinkedLists are also much less complicated than dealing with dynamic resizing, so that is another advantage, especially in an interview

**Problem:** *Maintain Linked List Tail Pointer.* head and tail are global pointers to the first and last elements, respectively, of a singly linked list of integers. Implement functions to delete and insert\_after. The argument to delete is the element to be deleted. The two arguments to insert\_after give the data for the new element and the element after which the new element is to be inserted. It should be possible to insert at the beginning of the list by calling insert\_after with None as the element argument. These functions should return 1 if successful and 0 if unsuccessful.

## Solution

Not really much to talk about, let's just code.

```
class Element:
    def __init__(self, data):
        self.data = data
        self.next = None

def locate(element):
    current, prev = head, None
    while current.data != element.data:
        if current.next is None: # Assumes tail.next is None
            raise ValueError(f"{element} not in linked list")
        prev = current
        current = current.next
    return prev, current

def delete(element):
```

```
if element is head:
    head = head.next
    return 1

try:
    current, prev = locate(element)
except ValueError:
    return 0 # Unsuccessful

if current is tail:
    tail = prev

prev.next = current.next
del current
return 1

def insert_after(element, data):
    new_element = Element(data)

    if element is None: # Insert at beginning
        new_element.next = head
        head = new_element
        return 0
    elif element is tail: # Insert at end
        current = tail
        tail = new_element
    else: # Lookup element
        try:
            current, _ = locate(element)
        except ValueError:
            return 0 # Unsuccessful

    new_element.next = current.next
    current.next = new_element
    return 1
```

Notes based on book solution:

- Make sure to update head and tail global pointers if they are affected by the insertion or deletion
- Account for corner cases of the list being empty or exceeding some maximum length that will overflow memory

**Problem:** *Bugs in remove\_head.* Find and fix the bugs in the following function that is supposed to remove the head element from a singly linked list:

```
def remove_head(head):
    del head
    head = head.next
```

## Solution

The bug in this code is that head is deleted prior to the pointer to the next node being utilized to set the new head of the linked list. Therefore, the interpreter will raise a NameError on the variable "head" when the second line of the function is executed. To fix the bug, we need to remove the `del head` line and leave the removal of the old head from memory up to garbage collection, or utilize a temporary variable to point to head such that we can remove it from memory ourselves. If the individual elements in this list take up large chunks of memory, the second approach is preferred such that we can ensure the freed-up memory is available immediately.

```
def remove_head(head):  
    head = head.next  
  
# OR to ensure memory is immediately free  
  
def remove_head(head):  
    temp = head.next  
    del head  
    head = temp
```

Notes based on book solution:

- Probably good practice to check that the object passed to the function has the next attribute and really is a pointer to the head of a linked list

**Problem:** *Mth-to-Last Element of a Linked List.* Given a singly linked list, devise a time- and space-efficient algorithm to find the mth-to-last element of the list. Implement your algorithm, taking care to handle relevant error conditions. Define mth-to-last such that when  $m=0$ , the last element in the list is returned.

## Solution

The brute force approach is to iterate through the list twice, the first time counting the elements and the second time accessing the mth-to-last element. This is space-efficient but not quite time-efficient (we can probably get creative and do better than  $O(2n)$ ).

A slightly better solution would be to iterate through the list fully once, and load each element into a hashtable with the index as the key. This solution would only require one iteration through the list and one constant time lookup using the hashtable, making it  $O(n+1)$ , but now we run into the problem of being space-inefficient since the hashtable will require just as much memory as the original list. Plus inserting each element into the hashtable is  $n O(1)$  operations, so this is basically  $O(2n+1)$  even with the hashtable.

My proposal is to use two pointers simultaneously. Start the first pointer at head. If the end of the list is reached before we have gotten to the mth element, handle the error case that the list

doesn't have an *m*th-to-last element. If the list has enough elements, start the second pointer at head once the first pointer reaches the *m*th element. Then step through the list incrementing both pointers at the same time. Once the first pointer reaches the tail, return the element where the second pointer is pointing.

```
def get_mth_to_last(head, m):
    if m < 0:
        raise ValueError("Value of m must be nonnegative")

    lead_pointer = head
    if lead_pointer is None:
        raise ValueError("Linked List cannot be empty")

    counter = -1
    rear_pointer = None
    while lead_pointer.next != None: # Indicating we're at the end
        if rear_pointer is None:
            counter += 1
            if counter == m:
                rear_pointer = head
        lead_pointer = lead_pointer.next
        rear_pointer = rear_pointer.next

    if rear_pointer is None: # Indicating list was shorter than m
        raise IndexError("linked list must contain at least m elements")

    return rear_pointer.data
```

Let's do some basic tests. If the list has a length of 10 and we use *m*=0, then neither value error is raised, nor is the index error raised. We begin the while loop. *rear\_pointer* is None, so *counter* is incremented. *Counter* is equal to *m*, so *rear\_pointer* is set to head (which is the same as *lead\_pointer*). Thus when we get to the end of the list, we will return the last element. We should probably raise a value error if *m* is not an integer at the top so the *m*<0 line doesn't throw an error. We should probably also check that the node passed through head has the necessary attributes of a singly linked list node. If *m* = 2, *rear\_pointer* gets set to head when *lead\_pointer* is head.next.next, so we will return the second-to-last element when *lead\_pointer* reaches the tail.

Notes based on book solution:

- None

**Problem:** *List Flattening*. Start with a standard doubly linked list. Now imagine that in addition to next and previous pointers, each element has a child pointer, which may or may not point to a separate doubly linked list. These child lists may have a child of their own, which may have its own child, and so on. Flatten the list so that all the nodes appear in a single-level, doubly linked list. You are given the head and tail of the first level of the list. Each node is an instance of the following class:

```
class Node:
    def __init__(self, data):
        self.next = None
        self.prev = None
        self.child = None
        self.data = data
```

## Solution

I would first want to know whether the final order mattered, since that would affect the ordering of the output. Let's assume it doesn't and then adapt the algorithm if necessary.

The multi-level data structure described in the problem is also referred to as an n-ary tree. Particularly since each node only has one possible child, this data structure can be thought of as a binary tree.

```
class Node:
    def __init__(self, data):
        self.next = None    # same as "right" in a BT node
        self.prev = None    # same as "parent" in a BT node
        self.child = None   # same as "left" in a BT node
        self.data = data
```

Therefore, one way we can flatten the data structure is traversing the data structure using either a breadth-first or depth-first search (depending on how the desired ordering of the output) and appending each element to a new array each time. However, this approach is spatially inefficient since copying the elements into a new array requires double the space.

To flatten the list without requiring additional space, we'll need to adjust the pointers instead of copying the elements. We can do this flattening slightly more efficiently by using a breadth-first search (ie. traverse the list with the next pointer, and every time a node has a child, load the child into a queue, then once you reach the tail pointer, set the next element to be the first element dequeued). This method only requires a lot of memory if the number of children in a particular level is huge, otherwise the space needed is much less than the full array.

The most efficient solution would not need a temporary data structure to store nodes in progress. Perhaps iterate through the list using the next pointer, and then whenever a node has a child, set the next pointer of tail to point to the child (and prev of child to point to tail), and then update the tail pointer to ... Hmm. Updating the tail pointer kind of sucks because you have to iterate through all the next attributes of the child. Perhaps a depth first search is best, and each time you encounter a child set it to be the next attribute. The issue here is losing the connection to the next attribute while processing the information from the child. I think the intermediate data structure might be the best we can do.

```
class Queue:
    def __init__(self):
```

```
self.queue = []

def enqueue(self, item):
    self.queue.append(item)

def dequeue(self):
    try:
        return self.queue.pop(0)
    except IndexError:
        raise IndexError("Queue is empty, nothing to dequeue")

def is_empty(self):
    return len(self.queue) == 0

def flatten_list(head, tail):
    if head is None or tail is None:
        raise ValueError("Cannot flatten empty list")

    queue = Queue()
    queue.enqueue(head)
    while not queue.is_empty():
        current = queue.dequeue()
        if current.child is not None:
            queue.enqueue(current.child)
        if current.next is not None:
            queue.enqueue(current.next)

    tail.next = current
    current.prev = tail
    tail = current
    del current.child
    del current.next
```

Notes based on book solution:

- Instead of loading into a queue, just stick the child at the end, and you'll eventually get to it to handle its children once your iteration reaches that point. This is  $O(n)$ , simpler, and requires no additional space for an intermediate data structure

```
def flatten_list(head, tail):
    if head is None or tail is None:
        raise ValueError("Cannot flatten empty list")

    current = head
    end_node = tail
    while current is not end_node:
        if current.child is not None:
            tail.next = current.child
            current.child.prev = tail
            current.child.next = None
            tail = current.child
```

```
current = current.next
```

**Problem:** *List Unflattening*. After having flattened the list above, restore the original multi-level data structure.

## Solution

We have the advantage that our approach never changed the child pointers of the nodes, so that is what we will have to use to reconstruct the multi-level list. Because we appended to the end, each time a node in the new list has a child, that indicates a new level. All we need to do is adjust the next / prev pointers and we should be good to go.

This will involve some recursion to explore the possible children of the child.

```
def unflatten_list(head, tail):
    if head is None or tail is None:
        raise ValueError("Cannot flatten empty list")

    def explore_and_separate(start):
        if start is None:
            return

        if start.child is None:
            return

        else:
            start.child.prev = None # Use None pointer to indicate new level
            start.child.prev.next = None
            explore_and_separate(start.child)

    current = head
    tail_updated = True # update tail after first child found, then use to terminate
    while current is not tail:
        if current.child is not None:
            if not tail_updated: # Update the tail when the first child is found
                tail = current.child.prev
                tail_updated = True

            explore_and_separate(current.child)

        current = current.next
```

Notes based on book solution:

- None

**Problem:** *Null or Cycle*. You are given a linked list that is either null-terminated (acyclic) or ends in a cycle (cyclic). Write a function that takes a pointer to the head of the list and determines if the list is cyclic or acyclic. You may not modify the list in any way. Note that the data contained



by a node might be identical to the data contained in another node, but that the two nodes are independent.

## Solution

Since we cannot modify the list, we are not able to add a boolean `has_visited` flag to each element as we traverse the list. This makes me think we should copy the elements to a secondary data structure. Let's think about a hashset (something with constant time lookup). We cannot use the data as the keys, since two nodes can have the same data. We could use a counter to index the nodes, but we risk updating the counter incorrectly when a cycle does occur. Hashset might be out.

Obviously our solution should have the property that if it encounters a pointer to `None`, it stops and returns "acyclic" as the answer. The brute force approach would copy every node as it is encountered and add a boolean flag, however this is spatially inefficient.

A fun python-specific approach is to store the memory address of each node in a constant-time lookup data structure, then ask if the next node in question's memory address is in the array. Since we'd only be storing memory addresses, the overhead would be significantly lighter than copying the whole list.

```
def is_cycle(head):
    if head is None:
        raise ValueError("Null pointer passed to is_cycle")

    memory_addresses = set()
    current = head
    while current is not None:
        if id(current) in memory_addresses:
            return True
        else:
            memory_addresses.add(id(current))
        current = current.next
    return False
```

Notes based on book solution:

- You could compare the node to all its predecessors by starting at the head and iterating until you reach the current node. If any nodes are equal then you've found a cycle. The trouble here is that this approach is  $O(n^2)$ .
- A clever solution is to use two pointers, and to increment them at different speeds. Therefore, if a cycle is present, the fast pointer will eventually equal the slow pointer. This approach is spatially efficient and  $O(n)$ .

```
def is_cycle(head):
    if head is None:
        raise ValueError("Null pointer passed to is_cycle")
```


```
fast_pointer = head
slow_pointer = head
try:
    while fast_pointer.next is not None:
        if fast_pointer is slow_pointer:
            return True
        fast_pointer = fast_pointer.next.next
        slow_pointer = slow_pointer.next

# Attribute error caused by end of list not having a next attribute, so acyclic
except AttributeError:
    return False

return False
```

## Chapter 4: Trees and Graphs

**Problem:** *Preorder Traversal.* Informally, a preorder traversal involves walking around the tree in a counter-clockwise manner starting at the root, sticking close to the edges, and printing out the nodes as you encounter them. For the tree to the right, the result is 100, 50, 25, 75, 150, 125, 110, 175. Perform a preorder traversal of a binary search tree, printing the value of each node.



### Solution

In a preorder traversal, we visit the root, then visit the left subtree, then visit the right subtree. We can straightforwardly implement this traversal with a recursive function.

```
class BSTNode:
    def __init__(self, data):
        self.left = None
        self.right = None
        self.data = data

    def __str__(self):
        return str(self.data)

def preorder_traversal(root):
    if root is None:
        return

    print(root)
    preorder_traversal(root.left)
    preorder_traversal(root.right)
```

Testing this solution on the example tree yields the output 100, 50, 25, 75, 150, 125, 110, 175 as intended.

Notes based on book solution:

- None

**Problem:** *Preorder Traversal, No Recursion.* Perform a preorder traversal of a binary search tree, printing the value of each node. This time, you may not use recursion.

### Solution

Sure. Why not. A preorder traversal of a binary tree is the same as a depth-first search. Thus, we can load the root onto a stack, and while the stack is not empty pop and print, then add the children to the stack. The only trick here is to match the preorder behavior, we have to load the right child onto the stack first, such that the left child is always last-in-first-out.

```
class Stack:
    # Stack.peek() is not needed for this example
    def __init__(self):
        self.stack = []

    def push(self, item):
        self.stack.append(item)

    def pop(self):
        return self.stack.pop()

    def is_empty(self):
        return len(self.stack) == 0

def preorder_traversal(root):
    stack = Stack()
    stack.push(root)
    while not stack.is_empty():
        current = stack.pop()
        print(current)
        if current.right is not None:
            stack.push(current.right)
        if current.left is not None:
            stack.push(current.left)
```

Testing on the example tree yields 100, 50, 25, 75, 150, 125, 110, 175 as intended.


Notes based on book solution:

- None

**Problem:** *Lowest Common Ancestor.* Given the value of two nodes in a binary search tree, find the lowest common ancestor. You may assume that both values already exist in the tree. For example, assume 4 and 14 are given as the two values for the tree to the right. The lowest common ancestor would be 8 because it's an ancestor to both 4 and 14 and there is no node lower on the tree that is an ancestor to both 4 and 14.

### Solution

Unfortunately we cannot treat the tree as an array because it is not complete. This has the feel of a recursive problem where for a given subtree you ask “are both 4 and 14 in my left (or right)



subtree?” and if the answer is yes you ask the same question of your left (or right) child. Once you get to a point where the answer is no, meaning one of 4 and 14 is in the left child and one is in the right child, you have found the lowest common ancestor. This approach is somewhat inefficient because it essentially does  $\log(n)$  depth first searches to verify that 4 and 14 are in the children, so maybe we can do better than reassessing the same nodes several times.

One property we did not use is the “search” property of the binary search tree (i.e. left child < root < right child). This may save us some time. Let’s start at 20. Since both 4 and 14 are less than 20, we know they both must be in the left subtree, so we go to 8. Since 4 is less than 8 and 14 is greater than 8, we know 4 must be in the left tree and 14 must be in the right tree. Thus, we’ve found our common ancestor by only visiting two nodes. This approach is  $O(\log(n))$ . Let’s implement it.

```
def find_lowest_common_ancestor(root, value1, value2):
    assert value1 < value2
    if value1 < root.data and value2 < root.data:
        find_lowest_common_ancestor(root.left, value1, value2)
    elif value1 > root.data and value2 > root.data:
        find_lowest_common_ancestor(root.right, value1, value2)
    elif value1 < root.data and root.data < value2:
        return root.data
    else:
        raise RuntimeError("Invalid binary search tree passed as root")
```

Testing the example, we start at 20 and we execute the if clause. We then move to 8, and this time we execute the second elif clause, returning the correct answer of 8.

Notes based on book solution:

- Since we’re not analyzing several branches and just going down the tree once, an iterative approach may be simpler than a recursive approach
- An iterative approach will also have fewer function calls, so in practice it may be faster than the recursive solution even though both are  $O(\log(n))$

## Chapter 5: Arrays and Strings

**Problem:** *First Non-repeated Character.* Write an efficient function to find the first non-repeated character in a string. For instance, the first non-repeated character in “total” is “o” and the first non-repeated character in “teeter” is “r”. Discuss the efficiency of your algorithm.

We'll definitely have to scan every character in the string, so efficiency-wise we are bounded by  $O(n)$ . Brute force would be to iterate through the string and for each character iterate through the remaining characters to determine whether the character appears again. This approach would be  $O(n^2)$ .

```
def find_first_nonrepeated_char(word):
    for idx, char in enumerate(word):
        repeated = False
        for ch in word[idx:]:
            if char == ch:
                repeated = True
                break
        if not repeated:
            return char
    # return a flag if no non-repeated characters are found
    return 0
```

Slightly more memory efficient would be to iterate through the string and track both the indices characters in a hash table. This would be  $O(n)$ , but would require additional space for the hash map. Although, the hashmap would only need as many keys as the number of unique characters in the word, so overall that's not very much additional space. In Python, we can use an `OrderedDict` to also handle keeping the entries in the order in which they appear in the word. Let's go with that.

```
from collections import OrderedDict

def find_first_nonrepeated_char(word):
    occur_dict = OrderedDict()
    for idx, char in enumerate(word):
        try:
            occur_dict[char] += 1
        except KeyError:
            occur_dict[char] = 1

    for char, occurrences in occur_dict.items():
        if occurrences == 1:
            return char
    # return a flag if no non-repeated characters are found
    return 0
```

Notes based on book solution:

- It would also be possible to use an array for the lookup, but this array will always be the size of the alphabet whereas the hashtable has the advantage of only needing to be the size of the number of unique characters in the word.

**Problem: Remove Specified Characters.** Write an efficient function that deletes characters from a string. For example, given a string of “Battle of the Vowels: Hawaii vs. Gronzy” and a removal argument of “aeiou”, the string should be altered to “Bttl f th Vwls: Hw vs. Grzny”. Justify any design decisions you make and discuss the efficiency of your solution.

So the brute for solution would be to iterate through the remove argument, and within that iteration, iterate through the string and remove any matches. This is  $O(n^2)$  and could certainly be improved.

I would load the remove argument into a hashtable, and then iterate through the string. If a string character is in the hashtable, then remove it from the string.

```
def remove_chars(string, remove):
    removal_map = {ch: None for ch in remove}
    out_string = []
    for ch in string:
        try:
            _ = removal_map[ch]
        except KeyError:
            out_string.append(ch)
    return ''.join(out_string)
```

In this solution, I'm choosing to append to a list and then join, rather than using string concatenation, because under the hood string concatenation with the “+” operator leads to a lot of unnecessary temporary variables. I'm also using a hash table for constant time lookup, and setting all the hash values to None to avoid unnecessary memory usage. The efficiency here is  $O(n + r)$  where  $n$  is the length of the string and  $r$  is the length of the removal argument.

Notes based on book solution:

- Building up a new string could become memory inefficient for long strings
- The proposal in the book suggests overwriting characters in the string as you iterate through it once. This works in C, but in Python strings do not support index-wise operations like insertion and deletion. Below is the book solution, which requires the extra step in Python of casting the string to a list

```
def remove_chars(string, remove):
    removal_map = {ch: None for ch in remove}
    string_arr = string.split('')
    source_counter = 0
    destination_counter = 0
    while source_counter < len(string_arr):
        try:
            _ = removal_map[string_arr[source_counter]]
```

```
except KeyError:
    string_arr[destination_counter] = ch
    destination_counter += 1
finally:
    source_counter += 1
return ''.join(string_arr[:destination_counter])
```

- Note that in Python, nothing is gained by this approach because we had to copy the string to an array, so we might as well have just used that memory to build up our string

**Problem: Reverse Words.** Write a function that reverses the order of the words in a string. For instance, your function should transform the string “Do or do not, there is no try.” to “try. no is there not, do or do”. Assume that all words are space delimited and treat punctuation the same as letters.

In Python, this is super easy and can be done using only standard built in functions.

```
def reverse_words(string):
    string_arr = string.split(' ')
    string_arr.reverse()
    return ' '.join(string_arr)
```

Let's try to solve it without built in functions too, just in case that is asked.

Since we'll have to scan the string for spaces, we'll definitely be at least  $O(n)$  for our approach. Since this problem was written for C, we'll use a list of characters for our string so that we have index-based operations available to us. The plan is to start a counter at the end of the string, iterate through the string backwards and decrement the counter as we go; every time we find a space, we put the chunk of the array between our counter and the end at the end of the array.

```
def reverse_words(string_arr):
    end = len(string_arr)
    back_counter = end
    counter = end - 1
    while counter >= 0:
        if string_arr[counter] == ' ':
            string_arr.extend(string_arr[counter:back_counter])
            back_counter = counter
            counter -= 1
    return string_arr[end:]
```

Notes based on book solution:

- A way to do this inplace is to iterate through the string twice. The first time, fully reverse it by swapping letters with their complements. Now the words will be in the correct order but the letters within the words will be backwards. The second time you iterate through, once you hit a space, stop, and swap all characters in that word, then continue. Technically this is still  $O(n)$  and does not require additional memory.



```
def reverse_string(string_arr, start, stop):
    while start < stop:
        string_arr[start], string_arr[stop] = string_arr[stop], string_arr[start]
        start += 1
        stop -= 1

def reverse_words(string_arr):
    reverse_string(string_arr, 0, len(string_arr))

    start_counter, stop_counter = 0, 0
    while stop_counter < len(string_arr):
        if string_arr[stop_counter] == ' ':
            reverse_string(string_arr, start_counter, stop_counter)
            start_counter = stop_counter
            stop_counter += 1
```

- Certainly more elegant, and creates the solution in place, but you need to make sure the string starts with a space such that the last word in the reversed string gets un-reversed

**Problem: Integer/String Conversions.** Write a function to convert an ASCII string to a signed integer, and write a function to convert a signed integer into an ascii string.

This question again feels weird for Python because builtin functions exist for type casting. The `int()` and `str()` functions already meet all the requirements of this problem. Regardless, let's pretend those functions don't exist.

To convert a string to an integer, we would want to first handle the sign, then we want to add the value of each digit to the result. This will be  $O(n)$ .

```
def str_to_int(string):
    if string[0] == '-':
        sign = -1
        end = 1
    elif string[0] == '+':
        sign = 1
        end = 1
    else:
        sign = 1
        end = 0

    digit_map = {'0': 0, '1': 1, '2': 2, ..., '9': 9}
    integer = 0
    power = 0
    start = len(string) - 1
    while start >= end:
        integer += digit_map[start] * 10**power
        power += 1
        start -= 1
```

```
    return sign * integer

def int_to_string(integer):
    sign = '-' if integer < 0 else ''
    exp = 1
    out_str = []
    zero = ord('0')
    abs_integer = abs(integer)
    while exp <= abs_integer:
        digit = (abs_integer // exp) % 10
        out_str.insert(0, chr(digit + zero))
        exp *= 10

    return sign + ''.join(out_str)
```

Both of these functions are  $O(n)$  where  $n$  is the number of digits in the integer / characters in the string. The only memory required in both cases is that of the final string / integer and the 10 key hashtable used for conversion.

Notes based on book solution:

- In the `str_to_int` function, you don't have to go backwards and track both the index and the exponent (the index already carries all the information that the exponent needs anyways). Instead, start from the front and everytime you move a place, multiply your number by 10. This optimization is called *Homer's Rule*.

```
def str_to_int(string):
    if string[0] == '-':
        sign = -1
        start = 1
    elif string[0] == '+':
        sign = 1
        start = 1
    else:
        sign = 1
        start = 0

    zero = ord('0')
    integer = 0
    for digit in string[start:]:
        integer *= 10
        integer += ord(digit) - zero
    return sign * integer
```

- I've revised the `int_to_str` function to use `ord` and `chr` functions instead of a hashtable, but the algorithm is the same as the solution

## **Chapter 6: Recursion**

## **Chapter 7: Other Programming Topics**

## **Chapter 8: Counting, Measuring, and Ordering Puzzles**

## **Chapter 9: Graphical and Spatial Puzzles**

## Online Google Interview Questions

From an online article: [The 15 most asked questions in a Google Interview](#)

**Problem:** *Delete Node with Given Key.* You are given the head of a linked list and a key. You have to delete the node that contains the given key.

### Solution

First I would like to know whether this is a singly linked list or a doubly linked list. I will proceed assuming singly linked, since it's slightly harder than deleting from a doubly linked list. The trick is connecting the previous node to the following node, so in a doubly linked list you have a pointer to the previous node, but in a singly linked list you don't have that luxury.

```
def delete_node_with_key(head, key):
    prev = None
    current = head
    while current.key != key:
        prev = current
        current = current.next

    prev.next = current.next
    del current
```

Notes based on online solution:

- Account for the possibility of the key not being in the list

**Problem:** *Copy Linked List with Arbitrary Pointer.* You are given a linked list where the node has two pointers. The first is the regular “next” pointer. The second pointer is called “arbitrary\_pointer” and it can point to any node in the linked list. Your job is to write code to make a deep copy of the given linked list. Here, deep copy means that any operations on the original list (inserting, modifying, or removing) should not affect the copied list.

### Solution

The issue here is that if you were to iterate through the list and copy nodes and pointers, you might get to the situation where the “arbitrary\_pointer” points to a node that has yet to be copied, so it would become a null pointer. To get around this, we might copy the values and next pointers first, and then copy the arbitrary pointers with a second iteration. This is  $O(n)$ , but requires visiting each node twice.

Come to think of it, copying next pointers kind of sucks too because they won't have anything to point to until the following node is copied. For the deep copy part of the problem, we'll need to create new nodes and set their pointers to other new nodes, so that will definitely be a part of the solution.

Another issue, setting the `arbitrary_pointer` requires you to have the node it's pointing to on hand, so even if the copy already exists, you still need an efficient way to find the copy. This lookup problem makes me want to use a hashmap to store the copies while their pointers are being set. Let's give that a try.

Here I'm going to assume the data values in the linked list are unique, such that I can use them as keys in a dictionary. If they are not unique, I would just use `id(node)` as the key in the dictionary. Actually, let's just do that for completeness.

I would also like to know about the data stored in the list nodes, to determine if making a deep copy of that data is an essential part of the problem. It seems like the point is to properly copy the pointers, so I will assume the data is strings or integers such that a copy and deep copy are the same.

```
class ListNode:
    def __init__(self, data):
        self.next = None
        self.arbitrary_pointer = None
        self.data = data

def copy_list_with_arbitrary_pointer(head):
    new_nodes = {}
    current = head
    while current is not None:
        key = id(current)
        # Copy the current node
        try:
            existing_copy = new_nodes[key]
        except KeyError:
            new_nodes[key] = ListNode(current.data)
            existing_copy = new_nodes[key]

        # Copy the next pointer
        if current.next is not None:
            next_key = id(current.next)
            try:
                existing_next = new_nodes[next_key]
            except KeyError:
                new_nodes[next_key] = ListNode(current.next.data)
                existing_next = new_nodes[next_key]
            existing_copy.next = existing_next

        # Copy the arbitrary pointer
        if current.arbitrary_pointer is not None:
            arbitrary_key = id(current.arbitrary_pointer)
            try:
                existing_arbitrary = new_nodes[arbitrary_key]
            except KeyError:
```



```
new_nodes[arbitrary_key] = ListNode(current.arbitrary_pointer.data)
existing_arbitrary = new_nodes[arbitrary_key]
existing_copy.arbitrary_pointer = existing_arbitrary

# Return the head of the copied linked list
return new_nodes[id(head)]
```

Here I have performed a deep copy by creating new `ListNode` instances for each node, setting the data attribute appropriately, and setting pointers only to copied nodes. This solution is also  $O(n)$  since all the nodes are only visited once, and since lookup in the dictionary is  $O(1)$ .

Notes based on online solution:

- Could just use the node itself in the hashtable instead of its id but the concept is the same

**Problem:** *Mirror Binary Trees*. Given the root node of a binary tree, swap the left and right children for each node.

### Solution

I swear if I get asked to invert a binary tree.... It's not hard, start at the root and recursively interchange the children of each node you visit.

```
def mirror_binary_tree(root):
    # Base condition
    if root is None:
        return

    # Swap
    root.left, root.right = root.right, root.left

    # Recurse
    mirror_binary_tree(root.left)
    mirror_binary_tree(root.right)
```

Notes based on online solution:

- My solution was preordered; theirs was postordered. I don't think it matters.

**Problem:** *Check if Two Binary Trees are Identical*. Given the roots of two binary trees, determine if these trees are identical or not.

### Solution

An  $O(n)$  solution here is to do either a depth first or breadth first traversal of both trees simultaneously, and check that the roots are always equal. No need to check the pointers to the children because if they are identical, the children get visited in the same order.

Actually, this doesn't account for the interchange of a left leaf with a right leaf for a node with only one child. The traversal order could be the same. We could either check the pointers to the children, but the simpler way is to do an inorder traversal, since we'll be able to tell if a child is left or right based on its ordering with respect to the root.

```
def are_identical(root1, root2):
    if root1 is None and root2 is not None:
        return False
    elif root2 is None and root1 is not None:
        return False
    elif root1 is None and root2 is None:
        return True

    left = are_identical(root1.left, root2.left)
    root = root1.data != root2.data
    right = are_identical(root1.right, root2.right)

    return left and root and right
```

Problem: Sum of two values

```
def sum_of_two_values(arr, sum_value):
    hashset = set()
    for value in arr:
        if sum_value - value in hashset:
            return True
        hashset.add(sum_value - value)

    return False
```

Problem: Move all zeros to the left.

```
def move_zeros_left(arr):
    # a la insertion_sort
    counter = 0
    while counter < len(arr):
        if arr[counter] == 0:
            arr.insert(0, arr[counter])
            counter += 1
```

I don't like it because it feels like every insertion can require at most n element shifts. There must be a better way

Problem: String Segmentation.

```
def segment(string):
    for idx in range(len(string)):
        first_word = string[:idx]
```

```
    if dictionary.contains(first_word):
        second_word = string[idx:]
        if (len(second_word) == 0 or
            dictionary.contains(second_word) or
            segment(second_word)):
            return True
    return False
```

Problem: Find all palindrome strings.

```
def find_palindromes(string, center_idx):
    l_counter, r_counter = center_idx, center_idx
    palindromes = []
    while l_counter >= 0 or r_counter < len(string):
        if string[l_counter] == string[r_counter]:
            palindromes.append(string[l_counter:r_counter+1])
            l_counter -= 1
            r_counter += 1
        else:
            return palindromes

def is_a_bunch_of_palindromes(string):
    l_counter = 0
    r_counter = 2
    all_palindromes = []
    while r_counter < len(string):
        if string[l_counter] == string[r_counter]:
            all_palindromes.extend(find_palindromes(string, l_counter + 1))

        l_counter += 1
        r_counter += 1

    return all_palindromes
```

Notes:

- You forgot even-length palindromes

## Extra Problems from Cracking the Coding Interview

<https://apprenticedeveloper.com/interview-coding-problem-strings-arrays/>

**Problem 1.1:** *Unique Characters*. Write an algorithm to determine if all the characters in a string are unique. What if you cannot use additional data structures.

### Solution

Easy way is to use a set and for each character add it to the set and ask whether it's in the set. This will be  $O(n)$  since lookup in the set will be  $O(1)$ . If possible, could we do it without the extra memory overhead? In the unique case, the set will contain all the characters, so we will require double the memory.

To handle the memory more efficiently, we could utilize the size of the ASCII alphabet to help us. Since there are only 256 characters, we could use an array of length 256 instead of a set. I would argue though, that if we know the size of the alphabet, no string with more than 256 characters will be unique, so the array method will always require a memory overhead of a length-256 array, but the set method will be more memory efficient if the string is not unique.

If we cannot use additional data structures, then we are forced to do element-wise comparisons, which will be an  $O(n^2)$  run time. If we are allowed to destroy the string, we can do a modified insertion sort in  $O(n \log(n))$  time (to avoid taking up extra space) and at each insertion check whether two characters are equal.

**Problem 1.2:** *String Reversal*. Write code to reverse a C-style string (string with a null character at the end). For Python, assume a new-line character.

### Solution

Here I propose to swap each element with its complement, perhaps using two counters, and stopping the iteration when the counters reach each other.

```
def reverse(string):
    l_counter = 0
    r_counter = len(string) - 2 # to avoid the '\n' character
    while l_counter < r_counter:
        string[r_counter], string[l_counter] = string[l_counter], string[r_counter]
        l_counter += 1
        r_counter -= 1
```

Now in Python, we cannot do index-based assignment in strings. Our only choice will be to create a new string. In this case, it's more memory efficient to append to a list instead of performing several string concatenations, since string concatenation requires temporary variables prior to assignment.

```
def reverse(string):
    out_str = []
    for ch in string[0:-1]:
        out_str.insert(0, ch)
    return ''.join(out_str) + string[-1]
```

**Problem 1.3:** *Remove Duplicate Characters.* Design an algorithm to remove duplicate characters in a string-array without using any additional data structures to copy the array.

### Solution

So here we're not allowed to use any sort of hashset to track the characters that have already appeared. We also cannot use something like counting sort where we track the occurrences of each. We'll have to go through the array and compare each character to all the others and do it in  $O(n^2)$  time. If we can jumble the returned string, we can do it a la modified insertion sort in  $O(n \log(n))$ .

If we're allowed a single variable of fixed memory size, I would use an array of length = `len(alphabet)` and count occurrences. Still, a set will outperform an array.

```
def remove_dup_chars(string_arr):
    letters = set()
    for idx, ch in enumerate(string_arr):
        if ch not in letters:
            letters.add(ch)
            string_arr[idx] = ''
    return ''.join(string_arr)
```

**Problem 1.4:** *Anagrams.* Write a method to determine if two strings are anagrams.

### Solution

To be anagrams, the words must contain all of the same letters. I would iterate through the first string and build up a hashtable with keys of the letters and values of the counts. Then I would go through the other string and decrease the counts in the hashtable. If I am left with any nonzero counts or the second array produces a key error, the strings are not anagrams.

```
def are_anagrams(string_1, string_2):
    ht = dict()
    for ch in string_1:
        if ch in ht:
            ht[ch] += 1
        else:
            ht[ch] = 0

    for ch in string_2:
```

```
try:
    ht[ch] -= 1
    if ht[ch] == 0:
        del ht[ch]
except KeyError:
    return False

return len(ht) == 0
```

If we can modify the strings (arrays) inplace and cannot utilize extra memory, insertion sort each one, and determine if they are equal for an  $O(n \log(n))$  solution.

**Problem 1.5: Replacement.** Write a method to replace all the spaces in a string (array) with '%20'.

Pythonically, this is straightforward: `string.replace(' ', '%20')`

If we want to solve the problem without the builtin method, while still treating the string as an array of characters, the problem is now to put in a string of length 3 where every space is. I guess we can do this linearly.

```
def replace(string_arr: list):
    counter = 0
    stop = len(string_arr)
    while counter < stop:
        if string_arr[counter] == ' ':
            string_arr.insert(counter, '%')
            string_arr.insert(counter + 1, '2')
            string_arr.insert(counter + 2, '0')
            counter += 2
            stop += 2
        counter += 1

    return string_arr
```

**Problem 1.6: Matrix Rotation.** Given an NxN matrix representation of an image, write a function to rotate the image by 90 degrees. Do the rotation inplace if possible.

### Solution

The non-inplace solution would be to copy each element to its rotated location in a new matrix. We can do that inplace, we just have to do a simultaneous four-element swap.

```
def rotate(image):
    levels = len(image) // 2
    counter = 0
    while counter < levels:
        (image[counter][counter], image[counter][counter-1],
```

```
image[counter-1][counter], image[counter-1][counter-1] =  
image[counter][counter-1], image[counter-1][counter-1],  
image[counter][counter], image[counter-1][counter])  
counter += 1
```

**Problem 1.7: Matrix Zeroing.** Write an algorithm such that if an element in a row or column is zero, its entire row and column are set to 0.

We don't want to iterate through the matrix if possible. This has the feel of some sort of weird matrix multiplication. In Python, we can use numpy to find the locations of all the zeros easily, but this will involve utilizing additional memory. Let's try to do it inplace.

In the brute force approach, we have  $O(n^2)$  where  $n$  is the dimensionality of the matrix, so since there are  $n^2$  elements in the matrix, one could argue that the brute force approach is linear with the number of elements.

The problem here is that the zeros we set will trigger the iteration later on, and will cause more rows / columns to be set to zero. We can get around this by giving up on inplace manipulation, since storing the locations of all the zeros (in the zero-matrix case) could be as large as the original matrix. We don't need the full locations though, just the unique row and column indices are enough.

```
def zero_matrix(arr):  
    rows = set()  
    cols = set()  
    for ii in range(len(arr)):  
        for jj in range(len(arr)):  
            if arr[ii][jj] == 0:  
                rows.add(ii)  
                cols.add(jj)  
  
    for ii in range(len(arr)):  
        for jj in range(len(arr)):  
            if ii in rows or jj in cols:  
                arr[ii][jj] = 0
```

**Problem 1.8: String Rotation.** Assume you have a method `isSubstring` which checks if one word is a substring of another. Given two strings, `s1` and `s2`, write code to check if `s2` is a rotation of `s1` using only one call to `isSubstring` (i.e., "waterbottle" is a rotation of "erbottlewat").

**Solution**